

Examen 1 - Exploration de donnes

Author: Ricardo Vallejo

Etape 1

1.1. Load Data

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_excel('USA_cars_dataset.xlsx')
df.head(5)
```

Out[1]:

	Unnamed: 0	price	brand	model	year	title_status	mileage	color	vin	lot	state	country	condition
0	0	6300	toyota	cruiser	2008	clean vehicle	274117.0	black	jtezu11f88k007763	159348797	new jersey	usa	10 days left
1	1	2899	ford	se	2011	clean vehicle	190552.0	silver	2fmdk3gc4bbb02217	166951262	tennessee	usa	6 days left
2	2	5350	dodge	mpv	2018	clean vehicle	39590.0	silver	3c4pdcgg5jt346413	167655728	georgia	usa	2 days left
3	3	25000	ford	door	2014	clean vehicle	64146.0	blue	1ftfw1et4efc23745	167753855	virginia	usa	22 hours left
4	4	27700	chevrolet	1500	2018	clean vehicle	6654.0	red	3gcpcrec2jg473991	167763266	florida	usa	22 hours left

1.2. Identifiez les différentes variables et leurs types. Résumez le tout dans un tableau.

```
In [2]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2499 entries, 0 to 2498
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Unnamed: 0            2499 non-null  int64  
1   price                 2499 non-null  int64  
2   brand                 2499 non-null  object  
3   model                 2499 non-null  object  
4   year                  2499 non-null  int64  
5   title_status          2499 non-null  object  
6   mileage               2499 non-null  float64 
7   color                 2499 non-null  object  
8   vin                   2499 non-null  object  
9   lot                   2499 non-null  int64  
10  state                  2499 non-null  object  
11  country                2499 non-null  object  
12  condition              2499 non-null  object  
dtypes: float64(1), int64(4), object(8)
memory usage: 253.9+ KB
```

TYPES VARIABLES ===== Integers variables: price year lot Float ou point dec variables: mileage Object structure: brand model year title status color vin state country condition


1.3. Déterminer le nombre de voitures vendues annuellement. Interprétez les résultats

In [3]: `df[['lot', 'year']].groupby('year').count()`

Out[3]:

	lot
year	
1973	1
1984	1
1993	1
1994	2
1995	1
1996	2
1997	2
1998	4
1999	1
2000	4
2001	5
2002	2
2003	9
2004	6
2005	6
2006	8
2007	6
2008	18
2009	11
2010	13
2011	23
2012	72
2013	86
2014	104
2015	196
2016	203
2017	377
2018	395
2019	892
2020	48

In [4]: `df[['lot', 'year']].groupby('year').count().plot(kind='bar')
plt.title('Cars vendus Histogram', color = 'black', fontsize = 16)
plt.xlabel('Years',color = 'black')
plt.ylabel('Number of Cars vendus', color = 'blue')
plt.xticks(color = 'blue')
plt.yticks(color = 'black')
plt.show()`



En el anne 2019 sont vendus 892 voitures, cest le anne de plus ventes entre 1973 et 2020. La plupart de vehicules sont vendus entre 2012 et 2019.

Étape 2 : Analyse des prix de voiture

1. À partir du fichier USA_cars_dataset.xlsx, créez une nouvelle structure df2010 qui contient les observations

entre 2010 (inclus) et 2020 (exlus).

In [5]: `yearscond = list(range(2009, 2020))`

In [6]:

is2010 = df['year'].isin(yearscond)

In [7]:

df2010 = df[is2010]

df2010.head(100)

Out[7]:

	Unnamed: 0	price	brand	model	year	title_status	mileage	color	vin	lot	state	country	condition
1	1	2899	ford	se	2011	clean vehicle	190552.0	silver	2fmdk3gc4bbb02217	166951262	tennessee	usa	6 days left
2	2	5350	dodge	mpv	2018	clean vehicle	39590.0	silver	3c4pdcgg5jt346413	167655728	georgia	usa	2 days left
3	3	25000	ford	door	2014	clean vehicle	64146.0	blue	1ftfw1et4efc23745	167753855	virginia	usa	22 hours left
4	4	27700	chevrolet	1500	2018	clean vehicle	6654.0	red	3gcpcrec2jg473991	167763266	florida	usa	22 hours left
5	5	5700	dodge	mpv	2018	clean vehicle	45561.0	white	2c4rdgeg9jr237989	167655771	texas	usa	2 days left
...
102	102	10780	ford	mpv	2017	clean vehicle	40455.0	white	1fm5k8ht0hga07252	167656360	texas	usa	2 days left
103	103	13800	ford	focus	2018	clean vehicle	23164.0	white	1fadp3j2xjl279400	167755491	south carolina	usa	21 hours left
104	104	25201	cadillac	door	2017	clean vehicle	19011.0	no_color	1gyknbrs8hz257399	167765111	michigan	usa	2 days left
105	105	7070	ford	mpv	2017	clean vehicle	45191.0	white	1fm5k7d82hgb39148	167656361	texas	usa	2 days left
106	106	8700	ford	focus	2018	clean vehicle	21405.0	white	1fadp3k23jl219764	167755494	south carolina	usa	21 hours left

100 rows × 13 columns

2. On s’intéresse uniquement à la variable price des voitures. Résumez dans un tableau deux indicateurs descriptifs de tendances centrales, deux de dispersions et deux de formes de cette variable.

Deux tendance central.

In [8]:

price = df2010['price']
price.head(5)

Out[8]:

1 2899
2 5350
3 25000
4 27700
5 5700
Name: price, dtype: int64

In [9]:

price.mode()

Out[9]:

0 16500
dtype: int64

In [10]:

price.mean()

Out[10]:

19189.131956155143

Deux dispersion

In [11]:

import statistics

statistics.stdev(price)

Out[11]:

11844.723105701421

In [12]:

statistics.variance(price)

Out[12]:

140297465.45073712

Deux de forme

In [13]:

price.skew()

Out[13]:

0.9656648298546364

3. Visualisez graphiquement un indicateur de tendance centrale et un indicateur de forme.

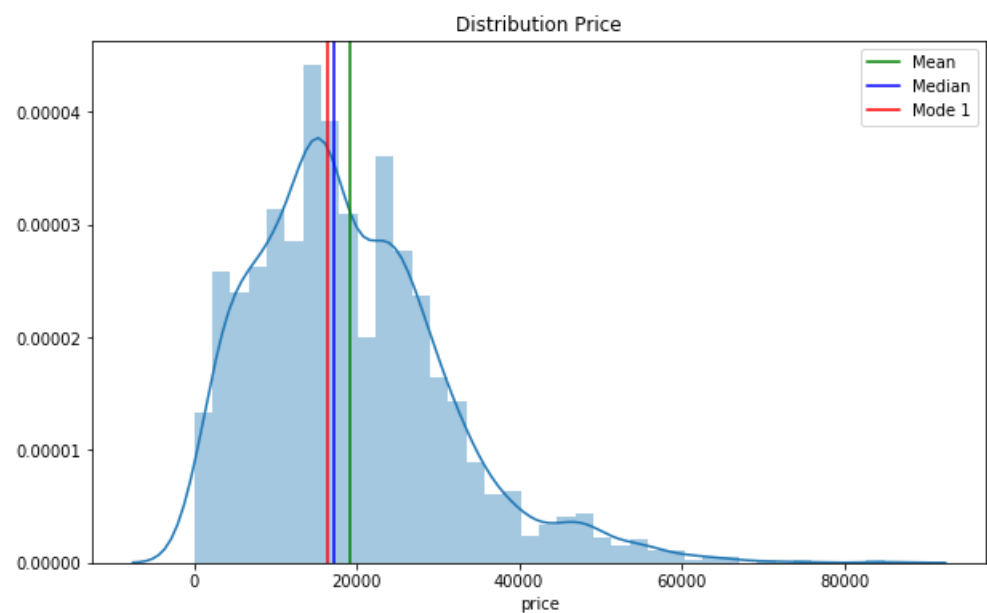
```
In [14]: import seaborn as sns

mean=price.mean();
median=price.median();
mode=price.mode();
fig, ax = plt.subplots(figsize=(10,6));

sns.distplot(price);
plt.title('Distribution Price');
plt.axvline(mean,color='green',label='Mean');
plt.axvline(median,color='blue',label='Median');
plt.axvline(mode[0],color='red',label='Mode 1')

plt.legend()
```

Out[14]: <matplotlib.legend.Legend at 0x1ea6d804e88>



Faites le lien entre les représentations graphiques et les indicateurs et commentez les résultats.

Le skew positive 0,96 nous confirme une distribution non simetrique. mean = 19189 Mode es 16500, ca veut dire que bcp de voitures sont vendus a cette price.

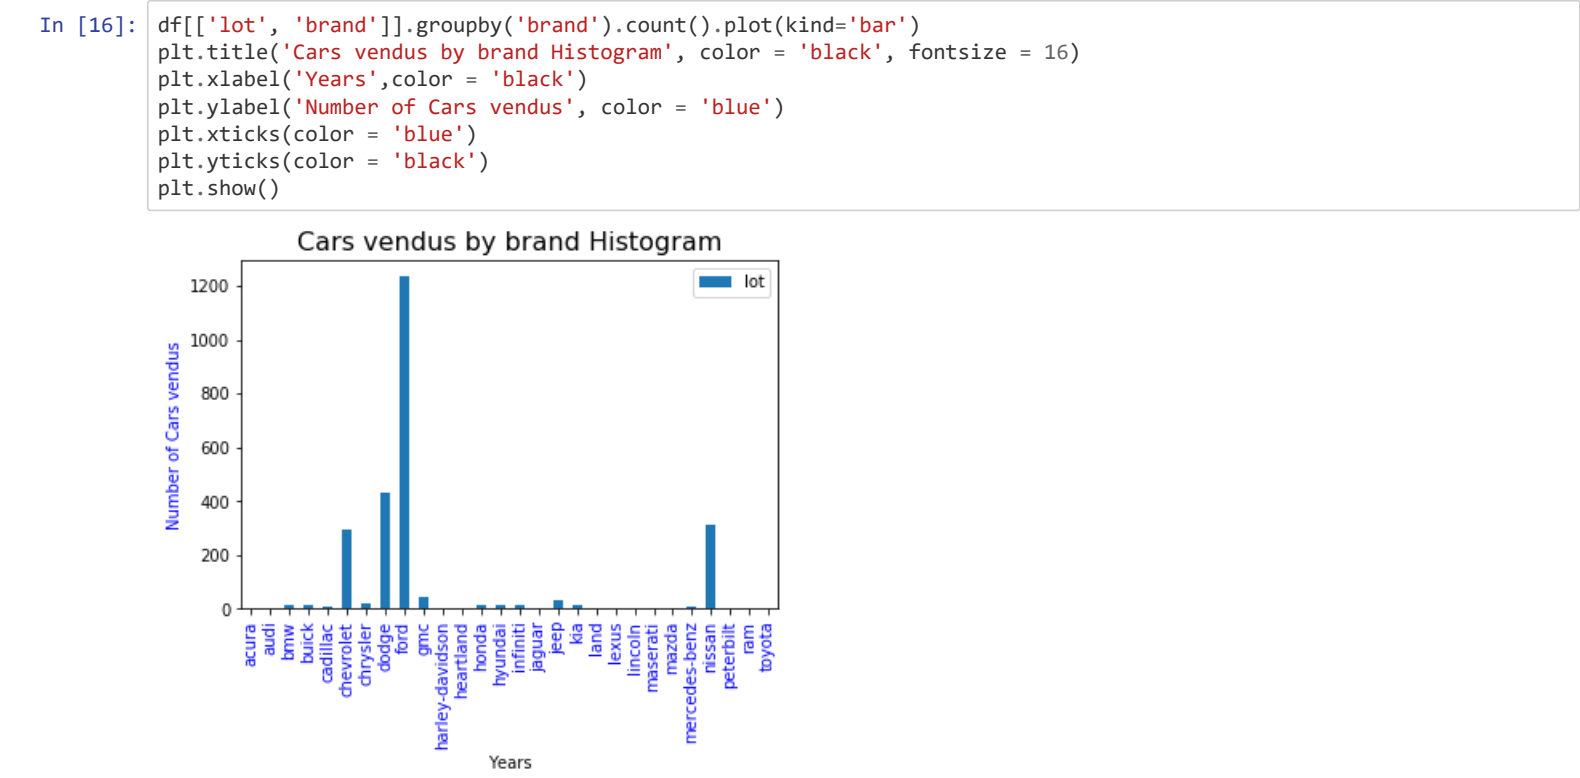
Étape 3 : Analyse des ventes de voiture

```
In [15]: dfventes = df2010[['lot', 'brand']].groupby('brand').count()
dfventes

dfventes.sort_values(by=['lot'], inplace=True, ascending=False)
dfventes2 = dfventes.head(6) #First 5 after order
dfventes2
```


Out[15]:

	lot
brand	
ford	1189
dodge	424
nissan	298
chevrolet	262
gmc	39
jeep	28



Les cars le plus vendus sont ford=1189, dodge=424 et chevrolet=262

2. Visualisez la variation de vente annuelle de chacune des 6 marques (sur un même graphique)



Étape 4 : Analyse des relations

Pensez-vous qu'il y a une relation entre le prix price et la distance parcourue mileage ? Expliquez la démarche et

les résultats obtenus.


```
In [18]: df2010.head(5)
```

Out[18]:

	Unnamed: 0	price	brand	model	year	title_status	mileage	color		vin	lot	state	country	condition
1	1	2899	ford	se	2011	clean vehicle	190552.0	silver	2fmdk3gc4bbb02217	166951262	tennessee	usa	6 days left	
2	2	5350	dodge	mpv	2018	clean vehicle	39590.0	silver	3c4pdcgg5jt346413	167655728	georgia	usa	2 days left	
3	3	25000	ford	door	2014	clean vehicle	64146.0	blue	1ftfw1et4efc23745	167753855	virginia	usa	22 hours left	
4	4	27700	chevrolet	1500	2018	clean vehicle	6654.0	red	3gcpcrec2jg473991	167763266	florida	usa	22 hours left	
5	5	5700	dodge	mpv	2018	clean vehicle	45561.0	white	2c4rdgeg9jr237989	167655771	texas	usa	2 days left	

```
In [19]: fig = plt.figure(figsize=(12,8))
plt.scatter( df2010['price'], df2010["mileage"], s=30)
plt.grid(color='gray', linestyle='--', linewidth=1)
plt.ylabel("Price")
plt.xlabel("Millaige")
plt.title('Variation de Price vs Millaige')
```

Out[19]: Text(0.5, 1.0, 'Variation de Price vs Millaige')



Graphiquement on a pas une relation lineaire entre les deux variables.

```
In [20]: # Regression Lineaire

from sklearn import linear_model
from sklearn.linear_model import LinearRegression

x1 = np.array(df2010['price']).reshape((-1,1))
y1 = df2010['mileage']

mymodel = LinearRegression()
results = mymodel.fit(x1,y1)

print("Coefficient determination: \n", results.score(x1, y1))
print("Intercept: \n", results.intercept_)
print("Slope: \n", results.coef_)
```

Coefficient determination:
0.12446583982016178
Intercept:
80379.37301767626
Slope:
[-1.62942198]

Le coefient de determination cest tres bas, donc pas de relation lineaire entre variables.

```
In [21]: from statsmodels.formula.api import ols

model2 = ols('price ~ mileage', data=df2010).fit() #x vs y
print(model2.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	price		R-squared:	0.124		
Model:	OLS		Adj. R-squared:	0.124		
Method:	Least Squares		F-statistic:	336.9		
Date:	Tue, 02 Mar 2021		Prob (F-statistic):	1.82e-70		
Time:	21:36:42		Log-Likelihood:	-25456.		
No. Observations:	2372		AIC:	5.092e+04		
Df Residuals:	2370		BIC:	5.093e+04		
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	2.294e+04	305.908	74.992	0.000	2.23e+04	2.35e+04
mileage	-0.0764	0.004	-18.355	0.000	-0.085	-0.068
=====						
Omnibus:	520.702	Durbin-Watson:	1.689			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1181.266			
Skew:	1.228	Prob(JB):	3.10e-257			
Kurtosis:	5.434	Cond. No.	9.88e+04			
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.88e+04. This might indicate that there are strong multicollinearity or other numerical problems.

R-squared confirm que on a pas relation linear entre variables price and mileage

3. Les variables prix price et title_status contiennent plusieurs valeurs aberrantes. En fixant un seuil à 2 (threshold = 2), supprimer les valeurs aberrantes en utilisant la distance interquartile.

```
In [22]: def is_outlier(value, p25, p75):
# Check if value is an outlier
lower = p25 - 1.5 * (p75 - p25)
upper = p75 + 1.5 * (p75 - p25)
return value <= lower or value >= upper

def get_indices_of_outliers(values):
#Get outlier indices (if any)
p25 = np.percentile(values, 25)
p75 = np.percentile(values, 75)

indices_of_outliers = []
for ind, value in enumerate(values):
if is_outlier(value, p25, p75):
indices_of_outliers.append(ind)
return indices_of_outliers

indices_of_outliers = get_indices_of_outliers(df2010['price'])
df2010['price'][indices_of_outliers] = mean; #np.percentile(df2010['price'], 75, interpolation = 'midpoint') - np.
percentile(df2010['price'], 25, interpolation = 'midpoint')
```

C:\Users\valm044\Anaconda3\lib\site-packages\ipykernel_launcher.py:21: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
-----
TypeError                                Traceback (most recent call last)
~\Anaconda3\lib\site-packages\pandas\core\series.py in __setitem__(self, key, value)
    1013         try:
-> 1014             self._set_with_engine(key, value)
    1015         except com.SettingWithCopyError:

~\Anaconda3\lib\site-packages\pandas\core\series.py in _set_with_engine(self, key, value)
    1053         try:
-> 1054             self.index._engine.set_value(values, key, value)
    1055         return

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.set_value()

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.set_value()

pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()

TypeError: '[38, 43, 88, 120, 252, 313, 326, 335, 341, 347, 422, 555, 576, 771, 775, 1016, 1119, 1161, 1196, 1199, 1228, 1231, 1232, 1235, 1237, 1238, 1240, 1241, 1242, 1243, 1244, 1269, 1270, 1286, 1296, 1315, 1324, 1347, 1351, 1354, 1359, 1373, 1461, 1536, 1538, 1597, 1600, 1670, 1749, 1752, 1786, 1789, 1791, 1867, 1935, 1941, 1945, 1947, 2083, 2085, 2087]' is an invalid key
```

During handling of the above exception, another exception occurred:

```
ValueError                                Traceback (most recent call last)
<ipython-input-22-fa86f6ac5400> in <module>
     19
     20 indices_of_outliers = get_indices_of_outliers(df2010['price'])
--> 21 df2010['price'][indices_of_outliers] = mean; #np.percentile(df2010['price'], 75, interpolation = 'midpoint') - np.percentile(df2010['price'], 25, interpolation = 'midpoint')

~\Anaconda3\lib\site-packages\pandas\core\series.py in __setitem__(self, key, value)
    1040         pass
    1041
-> 1042         self._set_with(key, value)
    1043
    1044         if cacher_needs Updating:

~\Anaconda3\lib\site-packages\pandas\core\series.py in _set_with(self, key, value)
    1090         if key_type == "integer":
    1091             if self.index.inferred_type == "integer":
-> 1092                 self._set_labels(key, value)
    1093             else:
    1094                 return self._set_values(key, value)

~\Anaconda3\lib\site-packages\pandas\core\series.py in _set_labels(self, key, value)
    1103         mask = indexer == -1
    1104         if mask.any():
-> 1105             raise ValueError(f"{key[mask]} not contained in the index")
    1106         self._set_values(indexer, value)
    1107

ValueError: [313 347] not contained in the index
```

Excercise 4


```
In [25]: from scipy import stats

results = stats.ttest_1samp(df2010['price'], 20000, 0)
print('stattiscit: ', results[0])
print('p_value: ', results[1])

stattiscit: -3.3341306580162486
p_value: 0.0008688293624366382
```

```
In [26]: # interpret p-value
alpha = 0.05
if results[1] <= alpha:
    print('(reject H0)')
else:
    print('(H0 holds true)')

(reject H0)
```

```
In [ ]: # Nos rejectons que la moyenne cest 20000
```